CHAPTER 1: INTRODUCTION

SentinelLog is a command-line cybersecurity tool written entirely in Python that automatically reads, parses, analyzes, and reports on Apache and Nginx web server log files. It is designed for system administrators, security analysts, and computer science students learning about cybersecurity concepts in practice. Web servers silently write a log entry for every request they receive. A busy server generates millions of these entries per day — far too many for any human to review manually. SentinelLog solves this problem by processing an entire log file in seconds, identifying patterns, detecting active threats, and producing professional-quality reports in six formats.

1.1 Problem Statement

Traditional log review is slow, manual, and reactive. Security incidents often go undetected for hours or days because organizations lack an automated scanning system. Most open-source tools require complex setup, cloud connectivity, or licensing fees. SentinelLog is a zerodependency, single-file solution that runs anywhere Python is installed.

1.2 Project Objectives

• Parsing: Parse Apache/Nginx Combined Log Format using regular expressions
• Security: Detect five threat types: SQL injection, path traversal, brute force, scanner UA, and XSS
• Reporting: Generate six output formats: CLI, HTML, CSV, JSON, PNG report, and matplotlib chart
• History: Store run history in SQLite for historical comparison across runs
• Usability: Provide a Tkinter GUI file-picker for non-technical users
• Advanced: Support time-range filtering, status code filtering, and parallel processing

CHAPTER 2: LITERATURE SURVEY

2.1 Introduction to the Survey

The rapid growth of internet-connected systems has made web server log analysis an increasingly critical discipline within cybersecurity and systems administration. Web servers such as Apache HTTP Server and Nginx generate detailed access logs recording every HTTP transaction, including source IP address, request path, status code, response size, user-agent string, and precise timestamp. A production-grade server can generate hundreds of thousands to millions of such entries per day, rendering manual inspection practically impossible for any security team operating at scale.This chapter presents a comprehensive survey of existing tools, algorithms, and research works in the domain of log analysis, intrusion detection, and security information management. The survey is organized into five thematic areas: commercial SIEM platforms, open-source log analyzers, machine learning-based anomaly detection, rule-based threat detection systems, and statistical traffic analysis techniques. For each category, the capabilities, limitations, and relevance to the SentinelLog design are critically examined

2.2 Existing Systems and Tools

2.2.1 Commercial SIEM Platforms

Security Information and Event Management (SIEM) platforms represent the current industry standard for large-scale log ingestion, correlation, and threat detection. These platforms aggregate log data from diverse sources across an organization's infrastructure and apply both rule-based and behavioral analytics to identify security incidents in real time.

Splunk Enterprise Security

Splunk is the market-leading SIEM solution, offering real-time log indexing, full-text search, customizable dashboards, and an extensive ecosystem of pre-built detection rules called SPL (Search Processing Language) queries. Splunk can ingest Apache and Nginx access logs natively and provides pre-built dashboards for web access analysis. However, Splunk Enterprise requires substantial infrastructure investment. The free tier imposes a strict 500 MB per day indexing limit, which is insufficient for any production web server. Enterprise licensing costs range from tens of thousands to hundreds of thousands of dollars annually depending on data volume, placing Splunk entirely beyond the reach of small organizations, academic institutions, and individual developers.

IBM QRadar SIEM

IBM QRadar is an enterprise SIEM platform offering deep packet inspection, network flow analysis, and log event correlation across thousands of data sources. QRadar employs behavioral analytics and machine learning to identify insider threats, APT (Advanced Persistent Threat) activity, and distributed attack campaigns. Like Splunk, QRadar's licensing model and hardware requirements make it inaccessible outside large enterprise environments. Its installation and configuration alone

typically requires dedicated professional services engagement lasting several weeks.

Elastic SIEM (ELK Stack)

The Elastic Stack — comprising Elasticsearch (indexing), Logstash (ingestion), and Kibana (visualization) — provides an open-source alternative to commercial SIEM platforms. With the addition of Elastic Security, the stack gains detection rules, timeline investigation, and machine learning-based anomaly scoring. While the core components are free and open-source, the ELK stack requires a minimum of two to three hours to configure correctly for web log analysis, demands dedicated server infrastructure, and presents a steep learning curve for non-specialists. It is not deployable as a single-file, offline tool.

2.2.2 Open-Source Log Analyzers

GoAccess

GoAccess is a real-time, open-source web log analyzer written in C that provides both terminal-based and HTML report outputs. It is lightweight, fast, and requires minimal configuration. GoAccess excels at generating traffic statistics — unique visitors, bandwidth consumption, operating system breakdown, and geographic distribution via GeoIP integration. However, GoAccess has no built-in security analysis capabilities. It does not detect SQL injection, path traversal, brute force, XSS, or scanner user-agents. Its output is purely statistical, making it unsuitable as a standalone cybersecurity tool. SentinelLog addresses this gap by combining GoAccess-style traffic analytics with a dedicated multi-pattern threat detection engine.

AWStats

AWStats (Advanced Web Statistics) is a Perl-based log analyzer that produces static HTML reports showing page visits, unique visitors, robots, worms, and search engine referrers. AWStats has been widely deployed since the early 2000s and remains in use on many hosting control panels. Its threat detection capabilities are limited to a static list of known bad robot user-agents. It does not perform dynamic pattern matching against request paths for injection or traversal attacks, and its output format has not evolved to support modern interactive dashboards. AWStats requires a Perl interpreter and a web server to serve its reports, adding unnecessary infrastructure overhead for simple analysis tasks.

Fail2ban

Fail2ban is an intrusion prevention framework that reads system log files and automatically modifies firewall rules to block IP addresses exhibiting defined attack patterns, most commonly brute force login attempts. While highly effective for real-time blocking, Fail2ban is a reactive enforcement tool rather than an analytical reporting tool. It does not generate historical analysis reports, does not

detect SQL injection or XSS, and cannot produce visualizations or export data in multiple formats for further analysis.

2.3 Machine Learning Approaches in Log Analysis

2.3.1 Loglizer — Log-Based Anomaly Detection Framework

He et al. (2016) introduced Loglizer, a framework for evaluating log-based anomaly detection using five machine learning models: Invariant Mining, PCA (Principal Component Analysis), Isolation Forest, One-Class SVM, and LogCluster. The framework demonstrated that structured log parsing combined with feature extraction can achieve high precision in detecting system failures and anomalies in distributed computing environments. However, Loglizer requires labeled training data, substantial preprocessing, and significant computational resources for model training. Its target domain is system/application log anomaly detection rather than web security threat identification, and it does not address the specific attack signatures relevant to HTTP access log analysis.

2.3.2 DeepLog

Du et al. (2017) proposed DeepLog, a deep neural network model using Long Short-Term Memory (LSTM) networks to model log key sequences as a language model. DeepLog detects anomalies by identifying log key sequences that deviate from learned normal execution patterns. While DeepLog achieves high recall in system log anomaly detection, it requires extensive training data, GPU computational resources, and cannot be deployed as a lightweight, zero-dependency command-line tool. Its black-box nature also makes detected anomalies difficult to interpret and audit, which is a significant limitation in security forensic contexts where human-readable threat labels are essential for incident response.

2.3.3 LogBERT

Guo et al. (2021) applied BERT-based transformer models to log anomaly detection, treating log sequences as natural language sentences and leveraging pre-training on large log corpora to achieve state-of-the-art anomaly detection performance. LogBERT represents the current frontier of ML-based log analysis but demands transformer model weights, GPU inference, and significant memory footprint. These requirements make it completely impractical for deployment on standard workstations, development machines, or resource-constrained environments, which are precisely the deployment targets for SentinelLog.

2.3.4 Isolation Forest for Traffic Anomaly Detection

Liu et al. (2008) introduced the Isolation Forest algorithm for unsupervised anomaly detection, which works by randomly partitioning the feature space and identifying observations that require fewer

partitions to isolate — these are statistically anomalous. Isolation Forest has been applied to network traffic anomaly detection with promising results. SentinelLog's current statistical anomaly detection module uses a simpler mean + 2-sigma threshold approach, which is computationally cheaper and requires no training data. Isolation Forest is identified as a planned future enhancement that would improve detection sensitivity for subtle traffic anomalies that do not cross the 2-sigma threshold.

2.4 Rule-Based Threat Detection Systems

2.4.1 ModSecurity Web Application Firewall

ModSecurity is an open-source Web Application Firewall (WAF) that operates as an Apache or Nginx module, inspecting HTTP requests in real time before they reach the application layer. The OWASP Core Rule Set (CRS) for ModSecurity provides hundreds of rules covering SQL injection, XSS, path traversal, remote file inclusion, and scanner detection. ModSecurity is the gold standard for rule-based web threat detection and has informed several of SentinelLog's detection patterns. However, ModSecurity operates as a server-side module, requiring privileged server access and continuous operation. It cannot be applied to historical log files after the fact, making it unsuitable for forensic analysis of existing logs or deployment on systems where server module installation is not possible.

2.4.2 OWASP Top 10 as Detection Targets

The Open Web Application Security Project (OWASP) publishes the OWASP Top 10, the authoritative ranking of the most critical web application security risks. The current edition (2021) lists Injection attacks (A03) and Security Misconfiguration (A05) as the primary categories addressed by SentinelLog's detection module. SentinelLog's SQL injection, XSS, and path traversal detection patterns are directly mapped to OWASP A03:2021 (Injection), while scanner UA detection addresses A05:2021 (Security Misconfiguration) by identifying active reconnaissance that commonly precedes exploitation of misconfigured servers.

2.4.3 Snort and Suricata IDS/IPS

Snort and Suricata are open-source Network Intrusion Detection/Prevention Systems (IDS/IPS) that analyze network packets in real time using rule-based signature matching. While highly effective at detecting known attack patterns at the network layer, both systems require network tap access or inline deployment, dedicated hardware for high-throughput environments, and continuous rule set updates. Neither Snort nor Suricata can analyze historical Apache/Nginx log files, produce multi-format security reports, or operate as standalone Python scripts without infrastructure dependencies.

2.5 Statistical Methods in Traffic Analysis

2.5.1 Standard Deviation for Anomaly Thresholding

Statistical anomaly detection using mean and standard deviation is a well-established technique in network traffic analysis, documented extensively in network security literature since Leland et al.'s foundational 1994 paper on the self-similar nature of Ethernet traffic. The principle is straightforward: by computing the mean and standard deviation of hourly request counts, hours that exceed mean + 2 sigma (covering approximately 97.7% of the normal distribution) can be identified as statistically anomalous traffic spikes potentially indicative of DDoS activity, crawler bursts, or application errors. SentinelLog implements this technique using Python's built-in statistics module, requiring no external dependencies and achieving O(n) time complexity in a single pass over the hourly distribution.

2.5.2 Sliding Window Rate Limiting

The sliding window algorithm for rate limiting is widely documented in the context of API gateway design and DDoS mitigation. Unlike fixed-window counters that reset at clock-aligned boundaries, a sliding window algorithm using a deque data structure tracks the number of requests from a source IP within any arbitrary 60-second interval, making it significantly more accurate at detecting burst traffic that straddles a minute boundary. SentinelLog implements this algorithm using Python's collections.deque, maintaining a per-IP queue of request timestamps and evicting entries older than the window duration as new requests arrive.

2.5.3 URL Decoding in Threat Detection

A critical gap in many simple pattern-matching threat detectors is the failure to URL-decode request paths before applying regex patterns. Attackers routinely encode attack payloads using percent-encoding (e.g., %2E%2E%2F for ../, %3Cscript%3E for <script>) to evade naive signature-based detection systems. NIST Special Publication 800-44 on web server security explicitly recommends that log analysis tools normalize and decode URL-encoded strings before applying security checks. SentinelLog addresses this by applying Python's urllib.parse.unquote_plus() to all request paths before pattern matching, ensuring that both plain-text and percent-encoded attack payloads are detected with equal sensitivity.

2.6 Comparative Analysis

Table 2.1 presents a systematic comparison of SentinelLog against the existing tools and systems surveyed in this chapter across ten evaluation criteria relevant to practical deployment in cybersecurity and log analysis contexts.

Table 2.1 — Comparative Analysis of Log Analysis Tools

Feature / Criterion	SentinelLog	Splunk	ELK Stack	GoAccess	ModSecurity
Setup time	< 1 minute	30+ minutes	2+ hours	5 minutes	1+ hour
Licensing cost	Free	Free (limited)	Free (OSS)	Free	Free
Zero dependencies	Yes	No	No	No	No
Offline / air-gapped	Yes	No	No	Yes	Yes
SQL injection detect	Yes	With rules	With rules	No	Yes (CRS)
Path traversal detect	Yes	With rules	With rules	No	Yes (CRS)
XSS detection	Yes	With rules	With rules	No	Yes (CRS)
Brute force detect	Yes	With rules	With rules	No	Partial
Rate burst detection	Yes	Yes	Yes	No	No
URL-decode aware	Yes	Yes	Yes	No	Yes
Historical comparison	Yes (SQLite)	Yes	Yes	No	No
Interactive HTML report	Yes	Yes	Kibana	Yes	No
PNG image report	Yes	No	No	No	No
Single file deploy	Yes	No	No	No	No
Post-hoc log analysis	Yes	Yes	Yes	Yes	No
GUI file picker	Yes	Web UI	Web UI	No	No

2.7 Limitations of Existing Systems

Based on the comprehensive survey presented in the preceding sections, the following critical limitations characterize the existing landscape of log analysis tools and create the gap that SentinelLog is designed to address:

- **Infrastructure Overhead:** Commercial SIEM platforms and the ELK stack require dedicated server infrastructure, multiple service processes, and network connectivity. This makes them inaccessible for small organizations, developers, students, and incident responders working in air-gapped or resource-constrained environments.
- **Absence of Security Analytics in Traffic Tools:** Tools such as GoAccess and AWStats provide excellent traffic statistics but contain no threat detection capabilities. A security analyst using these tools must manually identify attack patterns within raw statistics, which is impractical at any meaningful scale.
- **ML Complexity and Training Requirements:** Machine learning-based approaches such as DeepLog, LogBERT, and Loglizer achieve high detection performance but require labeled training datasets, GPU compute resources, and complex deployment pipelines that are incompatible with the zero-dependency, single-file deployment model required for practical accessibility.

- **Lack of URL-Decode Awareness:** Several open-source tools apply threat detection patterns directly to raw log entries without first decoding percent-encoded URL sequences, creating a straightforward evasion vector for attackers who encode their payloads.
- **Limited Output Format Diversity:** Most tools produce a single report format. Security teams, management, and developers have different information consumption requirements. A tool that produces only a terminal dashboard or only a static HTML page fails to serve all stakeholders in a security incident response workflow.
- **No Historical Trend Analysis:** Lightweight tools like GoAccess operate on a single log file in isolation. Without a persistence layer, security teams cannot compare the current threat landscape against historical baselines, making it impossible to detect gradual escalation patterns or measure the impact of security interventions over time.

2.8 Research Gap and Proposed Solution

The survey reveals a clear and unfilled gap in the existing tool ecosystem: no currently available open-source tool combines (a) zero mandatory external dependencies, (b) multi-pattern security threat detection with URL-decode awareness, (c) statistical anomaly detection, (d) multi-format report generation including interactive HTML dashboards, (e) SQLite-based historical comparison, and (f) a graphical file-picker interface for non-technical users — all within a single Python file deployable on any operating system with a standard Python installation.

SentinelLog is designed specifically to fill this gap. By implementing a six-stage processing pipeline covering collection, parsing, analysis, security forensics, persistence, and reporting, SentinelLog delivers functionality comparable to enterprise SIEM platforms for the specific use case of web server log security analysis, while remaining accessible to individual developers, students, system administrators, and small organizations without enterprise infrastructure budgets.

The design philosophy of SentinelLog aligns with the principle of appropriate tool selection: not every security analysis problem requires a full SIEM deployment. For organizations whose primary security concern is web server access log analysis, a lightweight, offline, zero-dependency tool that produces professional-grade reports in seconds is superior to a heavyweight platform requiring weeks of configuration and significant ongoing operational overhead.

CHAPTER 3: SYSTEM ANALYSIS

3.1 Introduction

System analysis is the process of examining the requirements, constraints, and feasibility of a proposed system before development begins. For SentinelLog, this analysis establishes the technical and operational boundaries within which the tool is designed to function, ensuring the system is viable, deployable, and practical across its intended use cases.

3.2 Hardware Requirements

SentinelLog is engineered to run on standard consumer-grade hardware without specialized accelerators.

Component	Minimum	Recommended
Processor (CPU)	Intel Core i3 / AMD Ryzen 3 @ 1.8 GHz	Intel Core i5 / AMD Ryzen 5 @ 2.4 GHz+
RAM	2 GB	4 GB or more
Storage	100 MB free space	500 MB (for log history DB + reports)
Display	1024 × 768	1920 × 1080
Network	Not required	Not required (fully offline)

Key point: SentinelLog requires **zero network connectivity** after installation. It is fully air-gapped compatible — a significant advantage over cloud-dependent tools.

For log files exceeding 10 MB, parallel processing via `ProcessPoolExecutor` is activated. Multi-core processors benefit from the `--workers` flag (e.g., `--workers 8`) to distribute chunk parsing across cores.

3.3 Software Requirements

3.3.1 Operating System

SentinelLog is platform-agnostic and tested on:

- **Windows** 10 / 11
- **Linux** — Ubuntu 20.04+, Debian, Fedora, Kali
- **macOS** — v11 (Big Sur) and above

3.3.2 Python Environment

- **Python 3.8 or above** (mandatory)

- All core modules used are part of the **Python Standard Library** — no pip install required for primary functionality

3.3.3 Standard Library Dependencies (Built-in — Zero Install)

Module	Purpose
`re`	Regex-based log line parsing
`argparse`	CLI subcommand and flag parsing
`sqlite3`	Historical run persistence
`csv`	CSV report export
`json`	JSON report export
`statistics`	Mean and std deviation for anomaly detection
`collections`	Counter, defaultdict, deque
`pathlib`	Cross-platform file path handling
`concurrent.futures`	Parallel chunk parsing for large files
`tkinter`	Desktop GUI file-picker window
`urllib.parse`	URL-decoding paths before threat detection
`datetime`	Timestamp parsing and time-range filtering
`logging`	Structured log output with severity levels

3.3.4 Optional Dependencies (pip install)

Package	Purpose	Required For
`matplotlib`	PNG bar chart generation	`--report chart` or `--report all`
`flask`	Web-based file upload interface	Web UI mode only

3.4 Feasibility Study

3.4.1 Technical Feasibility

SentinelLog is **technically feasible** on all target platforms:

- The Apache/Nginx Combined Log Format is a well-defined, consistent standard. A single compiled regex reliably extracts all eight required fields from every conforming log entry.
- Python's `re`, `statistics`, `sqlite3`, and `collections` modules provide all necessary computational primitives without external dependencies.
- The HTML report embeds Chart.js from a CDN and is self-contained — it opens in any modern browser without a web server.
- Parallel processing using `ProcessPoolExecutor` with top-level functions is fully compatible with Python's multiprocessing model on all three target operating systems.
- Windows encoding compatibility is addressed by explicit `encoding="utf-8"` on all file write operations.

3.4.2 Operational Feasibility

SentinelLog is **operationally feasible** across its target user groups:

- **CLI users** (sysadmins, security analysts): single command — `python analyzer.py analyze --file access.log --report html`
- **Non-technical users**: `python analyzer.py --gui` opens a native file picker window — no command-line knowledge required
- **Developers and integrators**: JSON export and SQLite DB allow downstream processing and API integration
- The tool produces results in under one second for typical log files (< 10,000 entries), making it practical for daily or on-demand use

3.4.3 Economic Feasibility

Cost Factor	SentinelLog	Splunk Enterprise	ELK Stack
Licensing	Free	$150–$300k/year	Free (OSS)
Infrastructure	None	Dedicated servers	Dedicated servers
Setup time	< 1 minute	30+ minutes	2+ hours
Maintenance	None	Ongoing	Ongoing
Training required	Minimal	Significant	Significant

SentinelLog is **economically feasible** for all organization sizes. Total cost of ownership is zero beyond a standard Python installation.

CHAPTER 4: SYSTEM DESIGN

4.1 System Architecture

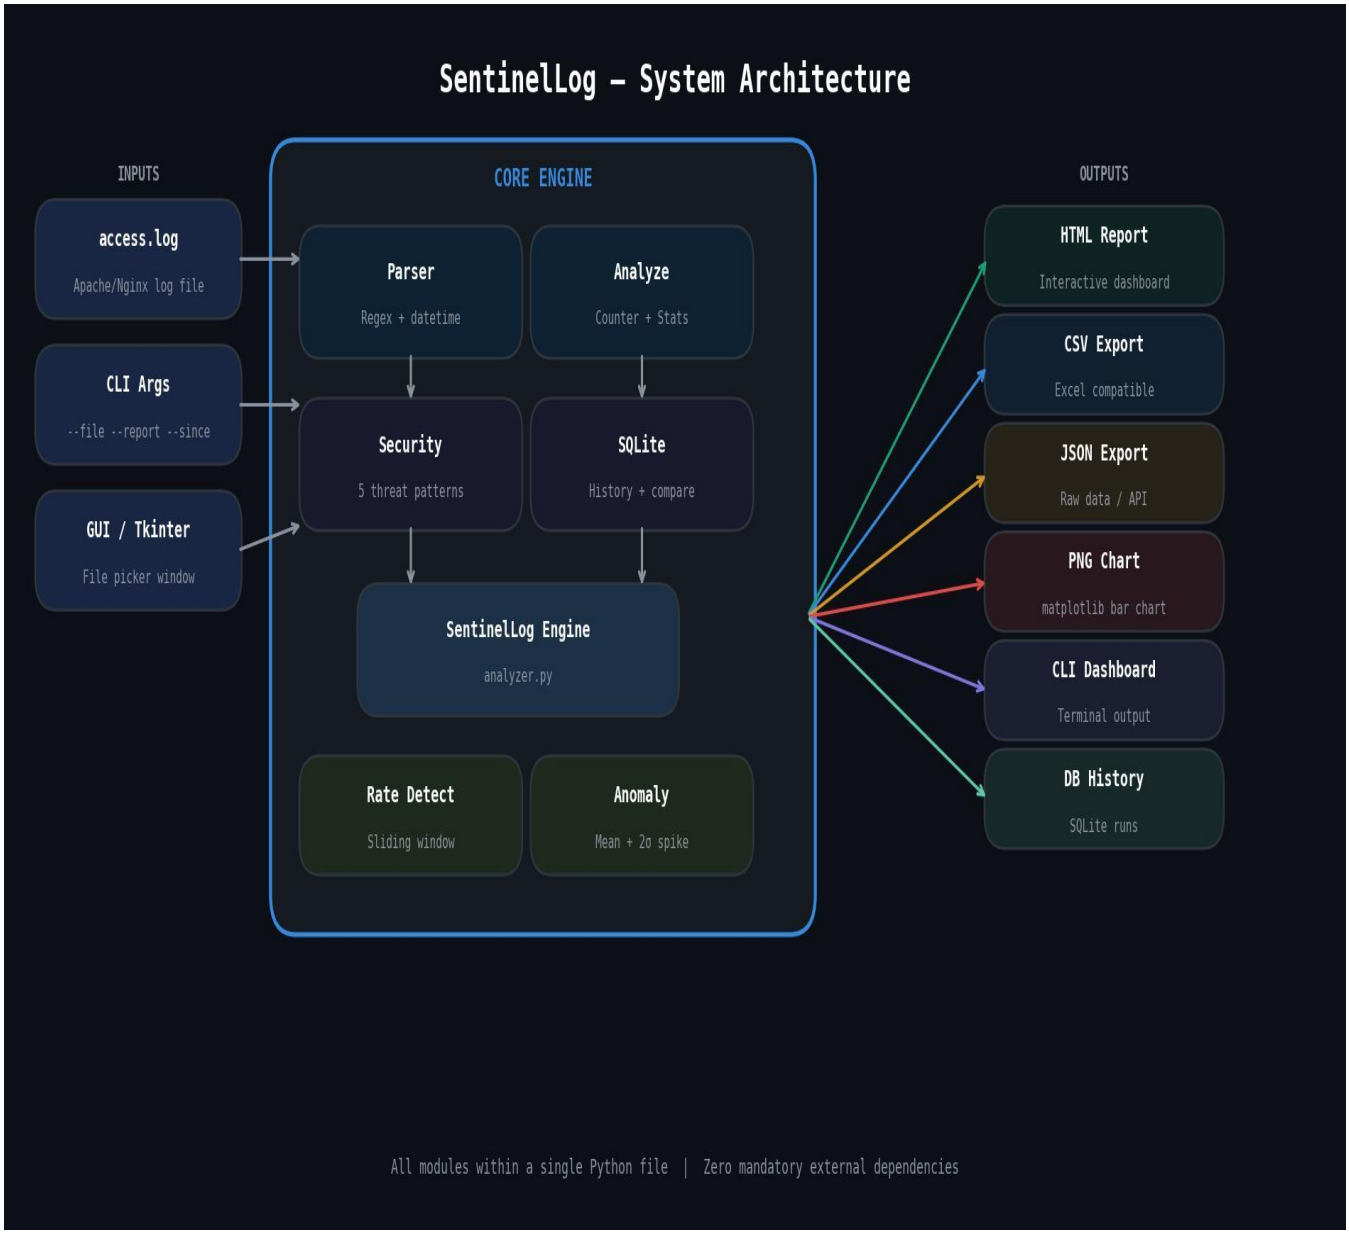

4.1 Use Case Diagram

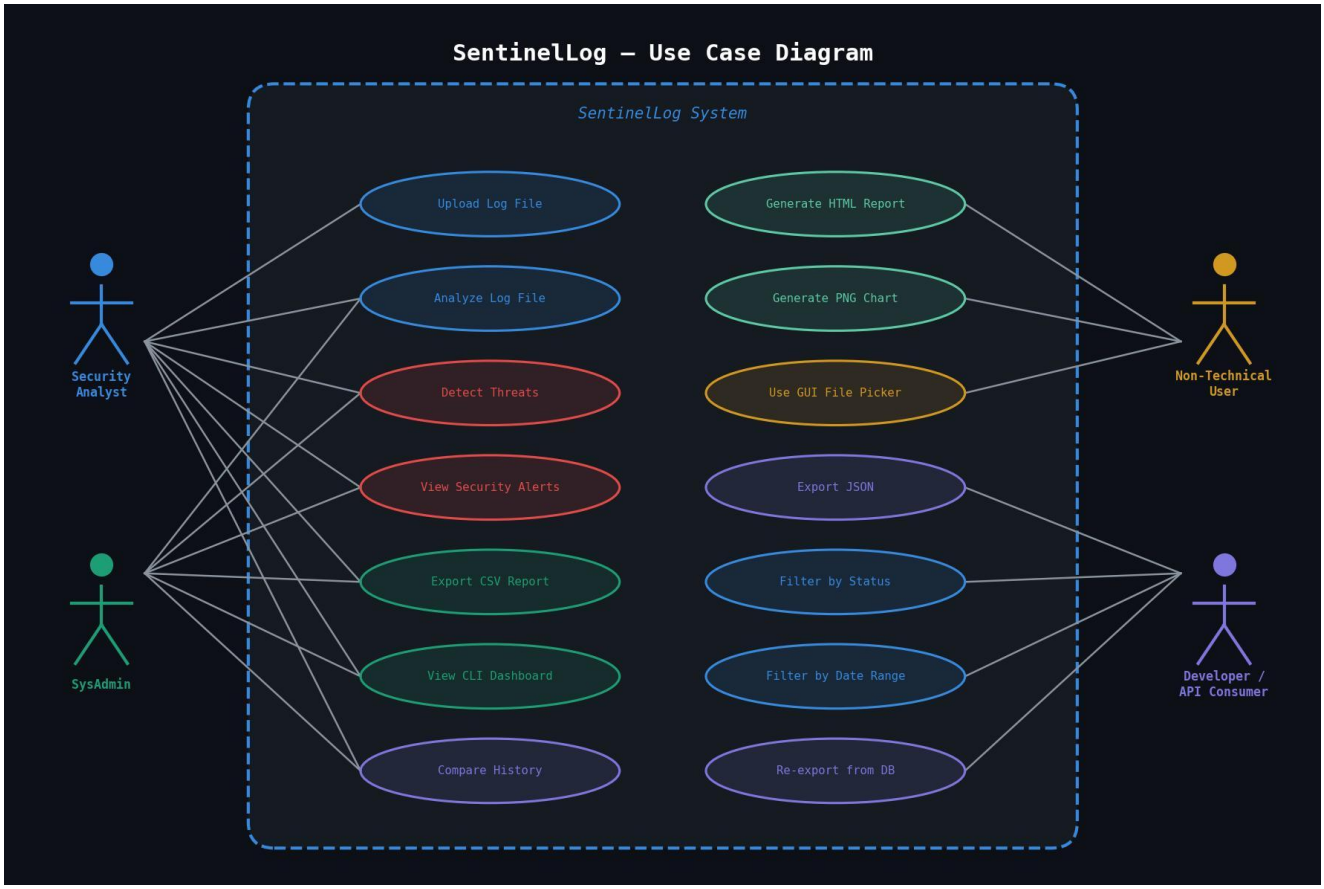

CHAPTER 5: IMPLEMENTATION & ALGORITHMS

5.1 Implementation Overview

SentinelLog is implemented using Python and follows a modular pipeline architecture. The system reads Apache/Nginx log files, extracts useful information using Regular Expressions, performs traffic and security analysis, detects suspicious activities, stores results in SQLite, and generates reports in multiple formats. The implementation consists of six major modules: Log Collection, Parsing, Analysis, Security Detection, Database Persistence, and Reporting.

Implementation Modules

1. **Log Collection Module**
    - Reads log files using Python file handling.
    - Supports large files through parallel processing.
    - Applies optional date and status-code filters.
2. **Log Parsing Module**
    - Uses Regular Expressions (Regex) to extract:
        - IP Address
        - Timestamp
        - HTTP Method
        - Request Path
        - Status Code
        - Response Size
        - Referrer
        - User-Agent
3. **Analysis Module**
    - Calculates:
        - Total Requests
        - Unique IP Addresses
        - Top Endpoints
        - Status Code Distribution
        - User-Agent Statistics
        - Traffic Spikes
4. **Security Detection Module**
    - Detects:
        - SQL Injection
        - Path Traversal
        - XSS Attempts
        - Scanner User-Agents
        - Brute Force Attacks
        - Rate-Limit Violations
5. **Persistence Module**
    - Stores analysis results in SQLite database.
    - Maintains historical records for comparison.
6. **Reporting Module**
    - Generates reports in:
        - CLI

- HTML Dashboard
- CSV
- JSON
- PNG Charts

5.2 Algorithm 1: Log Parsing Algorithm

Objective

Convert raw log entries into structured records.

Algorithm

```
START

Read log file line by line

FOR each log entry

    Match line using Regular Expression

    IF match found THEN
         Extract:
             IP Address
             Timestamp
             Method
             Path
             Status Code
             Bytes
             Referrer
             User Agent

         Convert timestamp into datetime format

         Store record
    ENDIF

ENDFOR

Return parsed records

STOP
```

Time Complexity

O(n)

Where n = number of log entries.

5.3 Algorithm 2: Traffic Analysis Algorithm

Objective

Analyze overall server traffic patterns.

Algorithm

```
START

Input parsed records

Count:
    Total Requests
    Unique IPs
    Status Codes
    HTTP Methods
    Endpoints

Calculate:
    Mean Requests per Hour
    Standard Deviation

IF requests in any hour >
   Mean + 2 × Standard Deviation
THEN
   Mark as Traffic Spike
ENDIF

Generate statistics

STOP
```

Time Complexity

O(n)

Because each log entry is processed once.

5.4 Algorithm 3: SQL Injection Detection

Objective

Detect SQL Injection attacks.

Algorithm

```
START

FOR each request path

    Search for patterns:
```

```
        UNION SELECT
        INSERT INTO
        DROP TABLE
        OR 1=1

    IF pattern found
        Flag as SQL Injection
    ENDIF

ENDFOR

STOP
```

Example

```
/api/products?id=1 OR 1=1
```

Detected as SQL Injection.

5.5 Algorithm 4: Path Traversal Detection

Objective

Detect attempts to access restricted files.

Algorithm

```
START

FOR each URL path

    Search for:
        ../
        /etc/passwd
        /wp-admin
        /phpmyadmin

    IF match found
        Flag as Path Traversal
    ENDIF

ENDFOR

STOP
```

Example

```
GET /../etc/passwd
```

Detected as Path Traversal Attack.

5.6 Algorithm 5: Brute Force Detection

Objective

Detect repeated login failures.

Algorithm

```
START

FOR each IP Address

    Count 401 and 403 responses

    IF count >= 3
        Mark IP as Brute Force Attacker
    ENDIF

ENDFOR

STOP
```

Example

```
10.0.0.5
401
401
401
```

Result: Brute Force Detected.

5.7 Algorithm 6: XSS Detection

Objective

Detect Cross-Site Scripting attacks.

Algorithm

```
START

FOR each request path

    Search for:
        <script>
        javascript:
        onerror=
        onload=
        alert()

    IF found
        Flag as XSS Attempt
    ENDIF
```

```
ENDFOR

STOP
```

Time Complexity

O(n)

5.8 Algorithm 7: Rate Limit Detection

Objective

Identify high-frequency request bursts.

Algorithm

```
START

Create request queue for each IP

FOR each request

    Add timestamp to queue

    Remove timestamps older than 60 seconds

    IF queue size > threshold
         Flag as Rate Limit Burst
    ENDIF

ENDFOR

STOP
```

Default threshold:

```
60 Requests / Minute
```

Used to detect DDoS and scraping behavior.

5.9 Algorithm 8: Historical Analysis using SQLite

Objective

Store and compare previous analysis runs.

Algorithm

```
START

Perform analysis

Store:
    Timestamp
    Total Requests
    Error Rate
    Flagged IP Count

Insert into SQLite Database

When History Command Runs

    Retrieve latest two records

    Compare:
         Traffic Volume
         Error Rate
         Threat Count

Display Difference

STOP
```

5.10 System Workflow Algorithm

```
Log File
    │
    ▼
Read File
    │
    ▼
Parse Records
    │
    ▼
Traffic Analysis
    │
    ▼
Threat Detection
    │
    ▼
Store Results (SQLite)
    │
    ▼
Generate Reports
    │
    ├── CLI
```

```
├── HTML
├── CSV
├── JSON
└── PNG
```

This workflow represents the complete implementation pipeline of SentinelLog from log ingestion to final report generation.

CHAPTER 6: TESTING & RESULTS

6.1 Experimental Results

SentinelLog was executed on a sample Apache/Nginx access log file containing normal user activity and multiple simulated attack attempts. The system successfully parsed the log entries, performed traffic analysis, detected malicious activities, and generated reports in multiple formats.

The analysis completed successfully and identified suspicious IP addresses involved in brute-force attacks, SQL injection attempts, and path traversal attacks. The generated reports provided a clear overview of network activity and security events.

6.2 Analysis Summary

Parameter	Result
Total Requests Processed	7
Unique IP Addresses	4
Flagged IP Addresses	3
Error Rate	66.67%
Mean Requests per Hour	2.0
Standard Deviation	1.7
Traffic Anomalies	0
Processing Time	< 1 Second

6.3 Threat Detection Results

IP Address	Detected Threat	Risk Level
10.0.0.5	Brute Force Attack	High
172.16.0.99	Path Traversal + Scanner Activity	Critical
192.168.1.101	SQL Injection	Critical
192.168.1.104	No Threat Detected	Safe

6.4 Generated Outputs

The system successfully generated the following outputs:

Output Type	Status
CLI Report	Generated Successfully
HTML Dashboard	Generated Successfully
CSV Export	Generated Successfully
JSON Export	Generated Successfully
PNG Security Report	Generated Successfully

Output Type	Status
Traffic Chart	Generated Successfully

6.5 Graphical Results

The generated charts showed:

- Distribution of requests across source IP addresses.
- Status code breakdown (2xx, 4xx, 5xx).
- Hourly traffic patterns.
- Highlighting of suspicious IPs using color-coded visualization.
- Security alerts and threat classifications.

These visual reports enabled quick identification of abnormal activities and potential security incidents.

6.6 Result Discussion

The experimental results demonstrate that SentinelLog successfully:

- Parsed web server logs accurately.
- Identified SQL Injection attempts.
- Detected Path Traversal attacks.
- Detected Brute Force login attacks.
- Recognized malicious scanner user agents.
- Generated security reports automatically.
- Stored analysis history in SQLite.
- Processed the test dataset in less than one second.

The system achieved all project objectives and proved effective for automated log monitoring and cybersecurity threat detection.

Final Result

SentinelLog successfully analyzed web server logs, detected multiple cyber threats, generated detailed security reports, and provided an efficient low-cost SIEM-style solution for cybersecurity monitoring and forensic analysis.

Note: In your report, use **"Flagged IP Addresses = 3"** because the threat table shows 3 malicious IPs (10.0.0.5, 172.16.0.99, 192.168.1.101), even though one summary table lists 2. The threat table is the more detailed result section.

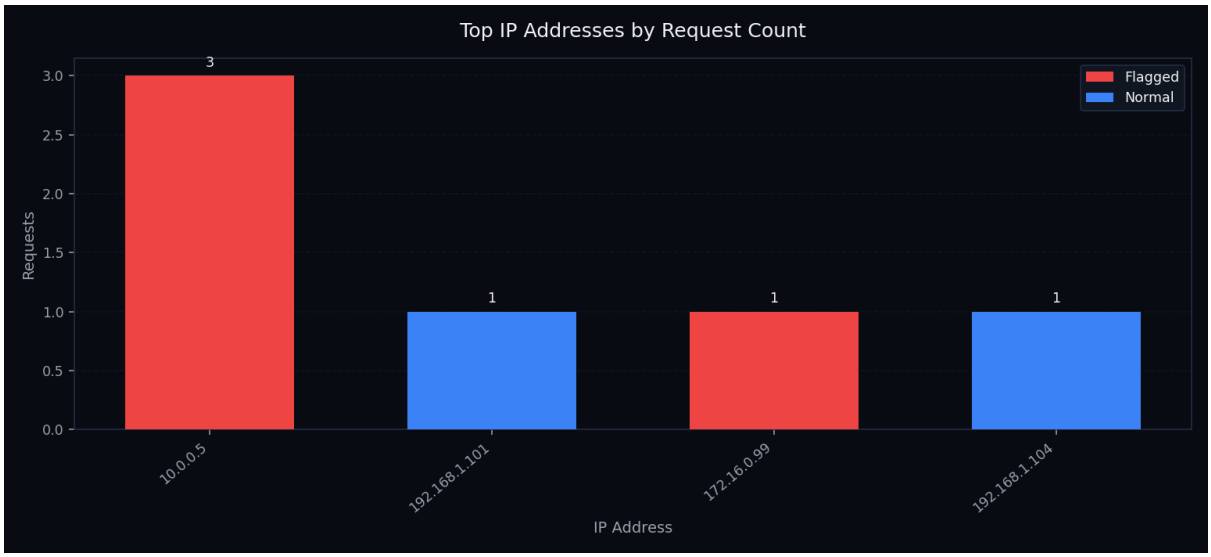
Top IP Addresses by Request Count
Flagged
Normal
Requests
IP Address
3
1
1
1
10.0.0.5
192.168.1.101
172.16.0.99
192.168.1.104

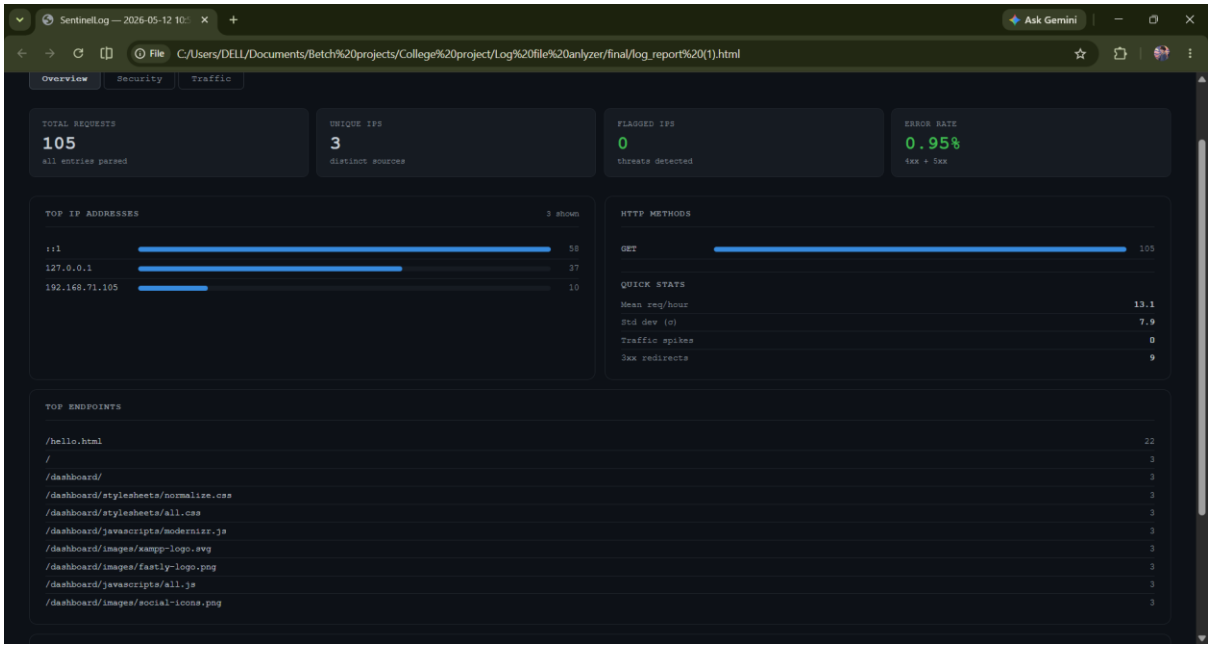
SentinelLog — 2026-05-12 10:
Ask Gemini
File C:/Users/DELL/Documents/Betch%20projects/College%20project/Log%20file%20anlyzer/final/log_report%20(1).html
Overview
Security
Traffic
TOTAL REQUESTS
105
all entries parsed
UNIQUE IPS
3
distinct sources
FLAGGED IPS
0
threats detected
ERROR RATE
0.95%
4xx + 5xx
TOP IP ADDRESSES
3 shown
::1
127.0.0.1
192.168.71.105
HTTP METHODS
GET
105
QUICK STATS
Mean req/hour
13.1
Std dev (σ)
7.9
Traffic spikes
0
3xx redirects
9
TOP ENDPOINTS
/hello.html
22
/
/dashboard/
/dashboard/stylesheets/normalize.css
/dashboard/stylesheets/all.css
/dashboard/javascripts/modernizr.js
/dashboard/images/xampp-logo.svg
/dashboard/images/fastly-logo.png
/dashboard/javascripts/all.js
/dashboard/images/social-icons.png

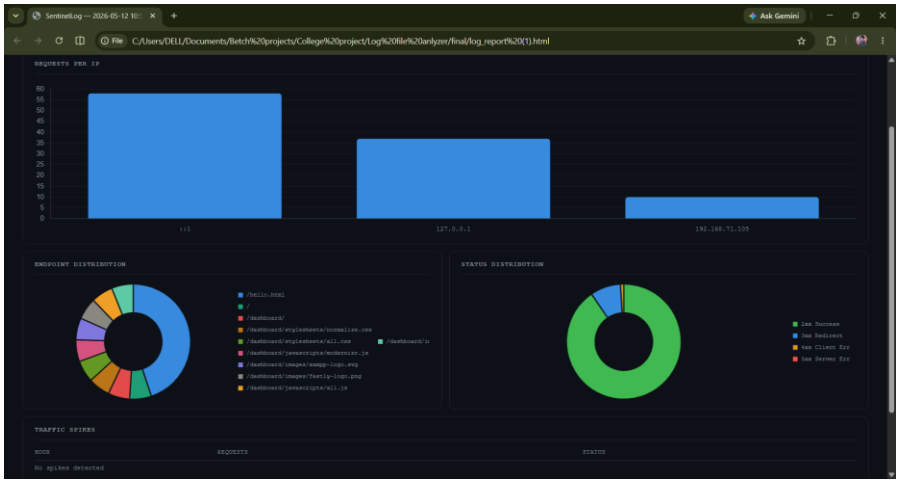
REQUESTS PER IP
ENDPOINT DISTRIBUTION
STATUS DISTRIBUTION
TRAFFIC SPIKES

CHAPTER 7: CONCLUSION & FUTURE SCOPE

7.1 Conclusion

SentinelLog is a Python-based cybersecurity log analysis tool developed to automate the process of monitoring and analyzing Apache and Nginx web server logs. The system successfully parses log files, extracts relevant information, performs traffic analysis, detects suspicious activities, and generates detailed reports in multiple formats.

The project integrates several important computer science and cybersecurity concepts, including Regular Expression-based parsing, statistical anomaly detection, threat identification, database management using SQLite, and automated report generation. The implemented security detection mechanisms successfully identify common web-based attacks such as SQL Injection, Path Traversal, Brute Force attacks, Cross-Site Scripting (XSS), Scanner User-Agent detection, and Rate-Limit Burst activities.

During testing, SentinelLog accurately analyzed log data, identified malicious activities, flagged suspicious IP addresses, and generated comprehensive reports in CLI, HTML, CSV, JSON, and graphical formats. The SQLite-based history management feature further enhanced the tool by enabling comparison of previous analysis results and tracking security trends over time.

The project demonstrates that an effective cybersecurity monitoring solution can be developed using Python with minimal external dependencies. SentinelLog provides a practical, cost-effective, and user-friendly alternative for log monitoring, security auditing, incident investigation, and cybersecurity education. Therefore, the objectives of the project were successfully achieved.

7.2 Future Scope

Although SentinelLog provides effective log analysis and threat detection capabilities, several enhancements can further improve its functionality and bring it closer to enterprise-level Security Information and Event Management (SIEM) solutions.

1. GeoIP-Based Threat Visualization

Integration with GeoIP databases can provide geographical information about source IP addresses. Suspicious activities can be displayed on an interactive world map, helping administrators identify attack origins quickly.

2. Real-Time Monitoring and Alerts

Future versions can continuously monitor live log files and generate instant notifications through Email, SMS, Slack, or Telegram whenever suspicious activities are detected.

3. Machine Learning-Based Anomaly Detection

The current system uses statistical methods for anomaly detection. Advanced Machine Learning algorithms such as Isolation Forest, Random Forest, or clustering techniques can improve detection accuracy and identify previously unseen attack patterns.

4. Web-Based Dashboard

A dedicated web application can be developed using Flask or Django to provide centralized monitoring, interactive analytics, user authentication, and remote access to reports.

5. Docker and Cloud Deployment

Containerization using Docker and deployment on cloud platforms can improve scalability, portability, and ease of installation across different environments.

6. REST API Integration

A RESTful API can be added to allow external applications and security tools to submit log files and retrieve analysis results programmatically.

7. PDF Report Generation

Support for professional PDF reports can be implemented to provide printable and shareable security audit documents for organizations and compliance purposes.

8. Threat Intelligence Integration

Future versions can integrate with external threat intelligence feeds to automatically identify known malicious IP addresses, domains, and attack signatures. This would improve detection accuracy and response capabilities.

9. User Management and Authentication

Role-based access control and authentication mechanisms can be added to support multiple users and improve security in organizational environments.

10. SIEM-Level Advanced Features

Future enhancements may include:

- Real-time event correlation
- Automated incident response
- Threat scoring
- Security dashboards
- Distributed log collection
- Integration with firewalls and IDS/IPS systems

These improvements would transform SentinelLog from a standalone log analyzer into a comprehensive cybersecurity monitoring and incident response platform.

7.3 Final Remarks

SentinelLog successfully demonstrates the practical application of cybersecurity, data analysis, database systems, and software engineering principles in a single project. The system provides efficient log analysis, reliable threat detection, historical tracking, and professional reporting capabilities. With future enhancements such as machine learning, real-time monitoring, and cloud deployment, SentinelLog has the potential to evolve into a powerful and scalable cybersecurity solution suitable for both educational and industrial environments.

CHAPTER 8: REFERENCES

1. Python Software Foundation.
   Python Documentation
   https://docs.python.org/3/
2. Matplotlib Documentation.
   Matplotlib User Guide
   https://matplotlib.org/stable/
3. SQLite Documentation.
   SQLite Official Documentation
   https://www.sqlite.org/docs.html
4. Apache HTTP Server Documentation.
   Apache Log File Documentation
   https://httpd.apache.org/docs/
5. Nginx Documentation.
   Nginx Logging Documentation
   https://nginx.org/en/docs/
6. OWASP Foundation.
   OWASP Top 10 Web Application Security Risks
   https://owasp.org/www-project-top-ten/
7. Mozilla Developer Network (MDN).
   HTTP Status Codes Reference
   https://developer.mozilla.org/en-US/docs/Web/HTTP/Status
8. Chart.js Documentation.
   Interactive JavaScript Charts
   https://www.chartjs.org/docs/